

# The AJPS Cryptosystem

Aggarwal, Joux, Prakash & Santha

Presented by Team HKV: Henry Herzfeld, Michael Keller, Yuri Villanueva

# Outline

- Motivation
- Background information: Mersenne primes, Hamming weights and distances
- The hard problem on which the cryptosystem is based
- How AJPS works
- Known attacks
- Our implementations
- Attacks that we tried
- Conclusions

# Motivation

- Simple
- Easier to understand
- Authors believe that it's secure against quantum attacks
- Previous attempt by MK unsuccessful

# Mersenne Primes

Mersenne number is

$$p = 2^n - 1$$

$$p = 0b11111\dots11111$$

when  $n$  is prime.

$p$  is a **Mersenne prime** when  $p$  itself is prime.

# The finite field $\mathbb{Z}_p$ when $p$ is a Mersenne prime

- $\mathbb{Z}_p$  is a finite field when  $p$  is prime.
- When  $p$  is a Mersenne prime,  $\mathbb{Z}_p$  has these nice properties:
  - $0b11111\dots11111 \equiv 0 \pmod{p}$
  - Multiplication by 2 is a circular bit shift

# Hamming Weight and Hamming Distance

- The **Hamming weight** of an integer  $m$ , written  $\text{Ham}(m)$ , is the number of 1s in its binary representation.
- The **Hamming distance** between two integers  $m$  and  $n$  is  $\text{Ham}(m \oplus n)$ .
- Hamming weight properties:
  - $\text{Ham}(x + y) \leq \text{Ham}(x) + \text{Ham}(y)$
  - $\text{Ham}(xy) \leq \text{Ham}(x) \cdot \text{Ham}(y)$
- Additionally, in  $\mathbb{Z}_p$ , with  $p = 2^n - 1$ , a Mersenne prime, we have the following nice properties:
  - For all  $i$ ,  $\text{Ham}(2^i x) = \text{Ham}(x)$ , because multiplication by  $2^i$  in  $\mathbb{Z}_p$  is just a cyclic bit shift.
  - $\text{Ham}(-x) = n - \text{Ham}(x)$  for all nonzero  $x$  in  $\mathbb{Z}_p$ . This is because  $0b11111\dots11111 \equiv 0 \pmod{p}$

# Security Assumptions

That is, on what hard problem is the AJPS cryptosystem based?

- Bit-by-bit- Mersenne Low Hamming Weight Problem
- The *Mersenne Low Hamming Combination Assumption*: Given an  $n$ -bit Mersenne prime and an integer  $h$ , such that  $4h^2 < n \leq 16h^2$ , the advantage of a probabilistic polynomial-time adversary in distinguishing

$$\left( \begin{bmatrix} R_1 \\ R_2 \end{bmatrix}, \begin{bmatrix} R_1 \\ R_2 \end{bmatrix} \cdot A + \begin{bmatrix} B_1 \\ B_2 \end{bmatrix} \right) \quad \text{and} \quad \left( \begin{bmatrix} R_1 \\ R_2 \end{bmatrix}, \begin{bmatrix} R_3 \\ R_4 \end{bmatrix} \right)$$

is at most  $O(2^{-h})$ , where  $R_1, R_2, R_3$ , and  $R_4$  are independent and uniformly chosen random  $n$ -bit strings, and  $A$  and  $B$  are independently-chosen  $n$ -bit strings of Hamming weight  $h$ .

# AJPS Single bit Encryption: Keygen

- Given the security parameter  $\lambda$ , choose a Mersenne prime  $p = 2^n - 1$  and an integer  $h$  such that  $\binom{n}{h} \geq 2^\lambda$  and  $4h^2 < n \leq 16h^2$ .
- Choose  $F, G$  to be two independent  $n$ -bit strings chosen uniformly at random from all  $n$ -bit strings of Hamming weight  $h$ .
- Set  $\text{pk} := H = \text{seq}\left(\frac{\text{int}(F)}{\text{int}(G)}\right)$ , and  $\text{sk} := G$ .



# AJPS Single Bit Encryption

**Encryption.** The encryption algorithm chooses two independent strings  $A, B$  uniformly at random from all strings with Hamming weight  $h$ . A bit  $b$  is encrypted as

$$C = \text{Enc}(\text{pk}, b) := (-1)^b (A \cdot H + B) \ .$$

# AJPS Single Bit Decryption

**Decryption.** The decryption algorithm computes  $d = \text{Ham}(C \cdot G)$ . If  $d \leq 2h^2$ , then output 0; if  $d \geq n - 2h^2$ , then output 1. Else output  $\perp$ .

For the correctness of the decryption note that  $C \cdot G = (-1)^b \cdot (A \cdot F + B \cdot G)$  which, has Hamming weight at most  $2h^2$  if  $b = 0$ , and at least  $n - 2h^2$  if  $b = 1$ .

Inefficient for large ciphertexts- multi bit variant required

# AJPS Multi-bit Encryption: Keygen

1. Choose a uniformly random  $n$ -bit integer  $R$  in  $\mathbb{Z}_p$ .
2. Choose random  $F$  and  $G$  among integers in  $\mathbb{Z}_p$  such that their binary representation has low Hamming weight  $h$ .
3. Compute  $T = F \cdot R + G \pmod{p}$
4. The public key is  $pk = (R, T)$ .
5. The private key is  $sk = F$ .

# AJPS Multi-bit Encryption: Encrypt

Encrypt with the public key  $(R, T)$

1. The scheme needs an efficient error-correcting code, with encoding function  $\mathcal{E} : \{0, 1\}^n \rightarrow \{0, 1\}^k$  and decoding function  $\mathcal{D} : \{0, 1\}^k \rightarrow \{0, 1\}^n$
2. To encrypt a message  $m \in \{0, 1\}^n$ , select three random numbers  $A$ ,  $B_1$ , and  $B_2$  of low Hamming weight  $m$ .
3. Output the ciphertext  $(C_1, C_2)$ , where

$$C_1 = A \cdot R + B_1 \text{ and}$$

$$C_2 = (A \cdot T + B_2) \oplus \mathcal{E}(m)$$

# AJPS Multi-bit Encryption: Decrypt

Decrypt  $C = (C_1, C_2)$  with the secret key  $F$  by computing the output  $\mathcal{D}(C_2 \oplus C_1 \cdot F)$

How does it work?

Recall that  $C = (C_1, C_2)$  with  $C_1 = A \cdot R + B_1$  and  $C_2 = (A \cdot T + B_2) \oplus \mathcal{E}(m)$ .

$$C_1 \cdot F = (A \cdot R + B_1)F = ARF + B_1 \cdot F = A \cdot T + B_1 \cdot F$$

Compare this with  $C_2 = (A \cdot T + B_2) \oplus \mathcal{E}(m)$ .

Since  $B_1$ ,  $B_2$ , and  $F$  have low Hamming weight  $h$ , the hamming distance between  $A \cdot T + B_1 F$  and  $A \cdot T + B_2$  is low, allowing them to almost cancel each other out with  $\oplus$ , leaving the a result that has low Hamming distance from  $\mathcal{E}(m)$ .

# Error-Correcting Code

- Error is random and spread out, so we can get away with a simple and efficient encoding.
- A simple repetition code works:  $k$  is the maximum message length in bits.  $\varrho$  is the number of repetitions of a single bit.
- Example: for message  $m = 0b1010$ ,  $\mathcal{E}(m) = 0b111...1000...0111...1000..0$ .
- The decoder  $\mathcal{D}(x)$  looks at blocks of  $\varrho$  bits and decodes each block to a single bit, 0 or 1 depending on which value is the majority in the block.
- The recommended parameters:  $n = 756,839$ ,  $k = 256$ , and  $\varrho = 2048$ .

# Random Oracle Implementation

```
m1, m2, m3 = {}, {}, {}
```

```
A = oracle(K, n, h, m1)  
B1 = oracle(K, n, h, m2)  
B2 = oracle(K, n, h, m3)
```

Memoization of returned pseudorandoms

Each memo (dictionary) defines a unique oracle

PRNG can also be constructed into oracle via explicitly defining seed prior to PR generation

```
def oracle(x, n, h, memo):  
    if x not in memo.keys():  
        memo[x] = get_nbit_ham_strings(n, h, 1).pop()  
    return memo[x]
```

# Key Encapsulation Mechanism (KEM)

Random Oracles

Why do we need it?

H\_1, H\_2, H\_3 are ROs

**Key Encapsulation.** Given the public key  $pk = (R, T)$ , the algorithm `Encaps` proceeds as follows:

1. Pick a uniformly random  $\lambda$ -bit string  $K$ .
2. Let  $A = \mathcal{H}_1(K)$ ,  $B_1 = \mathcal{H}_2(K)$ , and  $B_2 = \mathcal{H}_3(K)$ .
3. Let  $C = (C_1, C_2)$ , where  $C_1 = A \cdot R + B_1$ , and  $C_2 = \mathcal{E}(K) \oplus (A \cdot T + B_2)$ .
4. Output  $C, K$ .

**Decapsulation.** Given a ciphertext  $C = (C_1, C_2)$ , and  $sk = F$ , the decapsulation algorithm `Decaps` proceeds as follows:

1. Compute  $K' = \mathcal{D}((F \cdot C_1) \oplus C_2)$ .
2. Let  $A' = \mathcal{H}_1(K')$ ,  $B'_1 = \mathcal{H}_2(K')$ , and  $B'_2 = \mathcal{H}_3(K')$ .
3. Let  $C' = (C'_1, C'_2)$ , where  $C'_1 = A' \cdot R + B'_1$ , and  $C'_2 = \mathcal{E}(K') \oplus (A' \cdot T + B'_2)$ .
4. If  $C = C'$ , output  $K'$ , else output  $\perp$ .



# Previous Attacks

## N must be prime

Aggarwal et al. [1] propose several potential attacks on their system. We will summarize these here.

**Weak key attack.** Originally proposed by Beunardeau et al. [2], this attack is only relevant if all the active bits of  $F$  and  $G$  are in the less significant half of the string, that is, the right half of the string. If this is true, then  $F$  and  $G$  are smaller than  $\sqrt{P}$  and thus they can be easily recovered by a continued fraction expansion of  $H/P$ .

**LLL lattice attack.** Also proposed by Beunardeau et al. [2], this generalizes the weak key attack by guessing a decomposition of  $F$  and  $G$  into windows of bits so that all the active bits are on the right. By replacing the continued fraction method with LLL, it is possible to recover  $F$  from any possible window decomposition. Aggarwal et al. [1] revised their security parameter  $\lambda$  to 256 to counter this.

**Quantum attack using Grover's algorithm.** Using Grover's algorithm [4] for a quantum computer, one can speed up the lattice attack by a quadratic factor. Thus, we must ensure that our security parameter  $\lambda$  and  $h$  are equal.

**Meet in the middle attack.** This attack, proposed by de Boer et al. [3], has no effect on the security level of this cryptosystem for the chosen parameters, as its complexity is much larger than  $2^h$ .

**Active attacks.** Attacks of this type use the decryption of incorrectly encrypted ciphertexts to recover information about the secret key. To apply them to this cryptosystem, assume we have access to a decryption oracle. It is theoretically possible to leak information about the secret key by forming pseudo ciphertexts of the form  $A * H + B *$  with low Hamming weights that are not  $\lambda$ . To counter this, we can use the key encapsulation and decapsulation algorithms described in Section 1.

# Attack if $p$ is not Prime

If  $n_0 | n$ , then  $q = 2^{n_0} - 1$  divides  $p = 2^n - 1$ .  $F, G$  have Hamming weight  $\leq h \pmod q$ .

$Y = FR + G \pmod q$ . We can try to guess  $G$  from this, in time  $\sqrt{n_0}$  choose  $h$ .  
This will reveal  $F \pmod q$ , allowing us to guess  $F$  and  $G$  much faster than if  $p$  is prime.

# Our Implementations

We implemented AJPS three different ways:

- Representing large integers with sequences of bytes
- Using Python's built in large numbers, later replacing it with the GNU Multiprecision Arithmetic Library's ints.
- Representing low-Hamming-weight integers in  $\mathbb{Z}_p$  with a list of the positions of the 1s in their binary expansion.

```
# Repetition encoding. k is max bits in message, rho is number of repeated bits  
  
# k = maximum length of message in bits (256 bytes or UTF-8 characters)  
# rho = the number of repetitions of a bit
```

```
def E(m, k = 2048, rho = 256):  
    acc = 0  
    block_of_ones = 2**rho - 1 # 11111...111  
    while m:  
        if m & 1:  
            acc += block_of_ones  
            block_of_ones <<= rho  
            m >>= 1  
    return acc
```

```
def D(m, k = 2048, rho = 256):  
    acc = 0  
    block_of_ones = 2**rho - 1 # used as a mask for bitwise and  
    set_bit = 1  
    majority = (rho // 2) + 1  
    while m:  
        if hamming_weight(m & block_of_ones) >= majority:  
            acc |= set_bit  
            m >>= rho  
            set_bit <<= 1  
    return acc
```

# Using Python's built-in bignum and Bit Shifts

- At first, no special data structures were used--just Python's built-in bignums.
- Because of  $\mathbb{Z}_p$ 's special properties, we liberally used bit shifts and bitwise logical operators.
- It handles 756,839-bit numbers just fine.
- On a 2013 Mac Pro, it takes about 1.5s for keygen, and around 0.8s-1s for encrypting and decrypting a 44-byte message.
- Performance increases by over a factor of 10 with gmpy2.

```
def get_random_int_of_hamming_weight_h(h, n):  
    a = 0  
    for i in range(h):  
        bit_mask = 1 << randrange(n)  
        while a & bit_mask:  
            bit_mask = 1 << randrange(n)  
        a = a | bit_mask  
    return a
```

```
def hamming_weight(a):  
    acc = 0  
    while a:  
        if a & 1: acc += 1  
        a >>= 1  
    return acc
```

```
def hamming_distance(a, b):  
    return hamming_weight(a ^ b)
```

# Using Bit Position Lists to save space

- With the recommended parameters, the secret key  $F$  is an integer up to 756,839 bits long, but has Hamming weight 256. This takes up almost 100 kilobytes of space.
- Same goes for  $G$ ,  $A$ ,  $B_1$ , and  $B_2$ .
- If we represent these sparse integers as lists containing the positions of the 1s in their binary expansion, then each one takes up only  $256 \times 4$  bytes = 1 kilobyte.
- Even less storage is needed if we use 20 bits per position—5120 bits, comparable to the length of an RSA key..
- Implementation using these bit-position lists is somewhat slower than the one with Python bignums, within a factor of 2.
- Performance also improves dramatically with gmpy2.

```
def int_to_bit_position_list(a):  
    acc = []  
    i = 0  
    while a:  
        if a & 1: acc.append(i)  
        i += 1  
        a >>= 1  
    return acc
```



```
def int_plus_bpl(a, bpl):  
    return a + bit_position_list_to_int(bpl)
```

```
def int_times_bpl(a, bpl, pp = p):  
    bpl.sort()  
    acc = 0  
    i = 0  
    for bp in bpl:  
        a = (a << (bp - i)) % pp  
        acc = (acc + a) % pp  
        i = bp  
    return acc
```

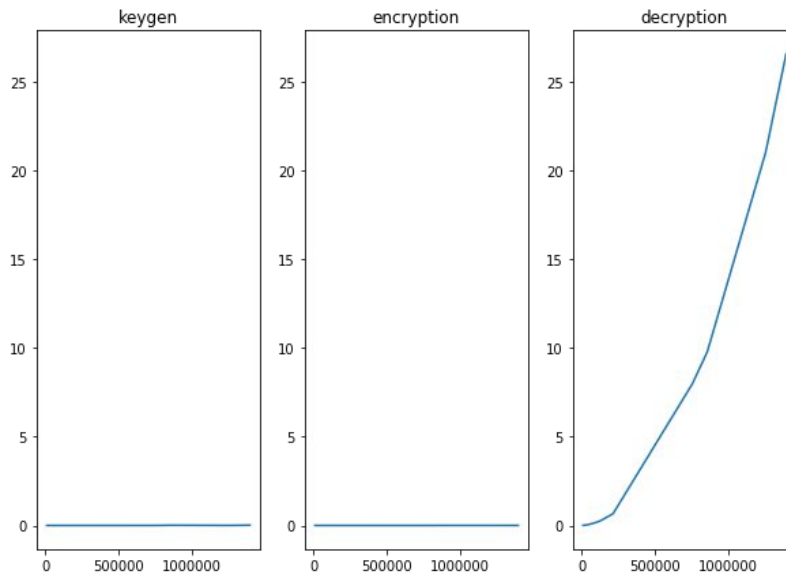
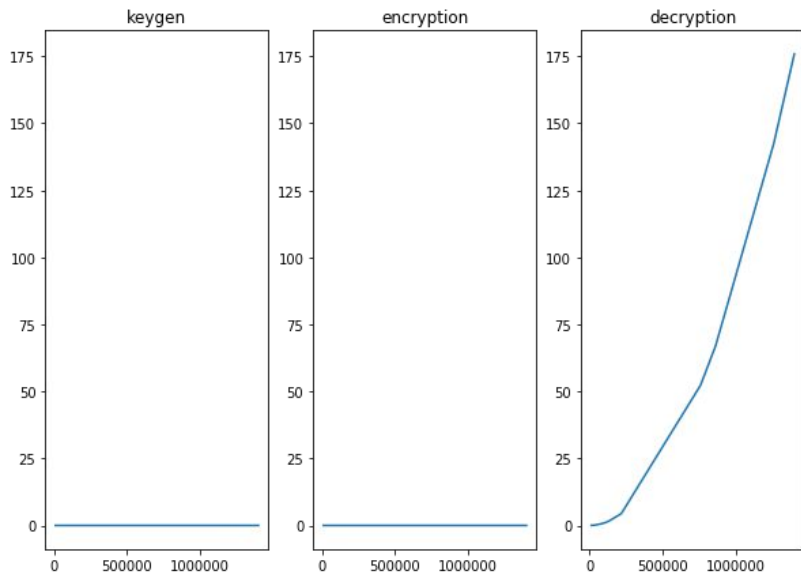
```
def bit_position_list_to_int(bpl):  
    return int_times_bpl(1, bpl)
```

```
def get_random_bpl_of_hamming_weight_h(h, n):  
    bpl = list(map(int, np.random.choice(n, h, replace = False)))  
    bpl.sort()  
    return bpl
```

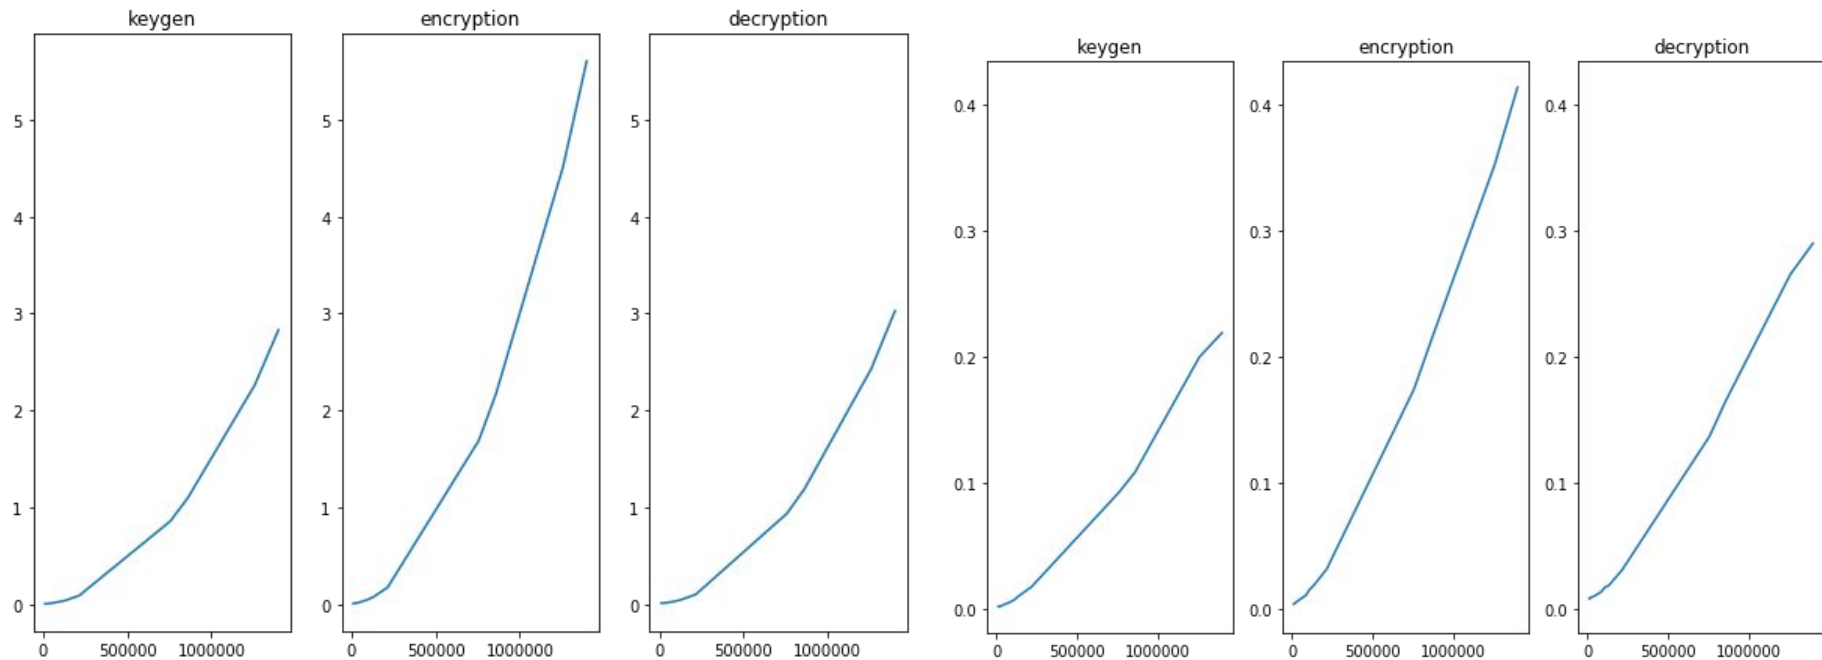
# AJPS Full



# AJPS bit-by-bit



# AJPS Full with Bit Position List





# Our attempts at attacking AJPS

- Brute Force
- Meet in the middle (MITM)
- Physical attacks on the low-Hamming-weight noise

# Brute Force

Search space is  $N$  choose  $W$

Running time:

Classical:  $O(n^{\text{srt}(n)/4})$

Quantum:  $O(n^{\text{sqrt}(n)/8})$

```
pk, sk = keygen(n, h)

guess = get_nbit_ham_strings(n, h, 1).pop()

while ham(guess//pk) != h or guess//pk != sk:
    guess = get_nbit_ham_strings(n, h, 1).pop()
```

Upon announcement of AJPS, team claimed this was the optimal solution.

# Meet-in-the-Middle Attack

Subset-sum problem

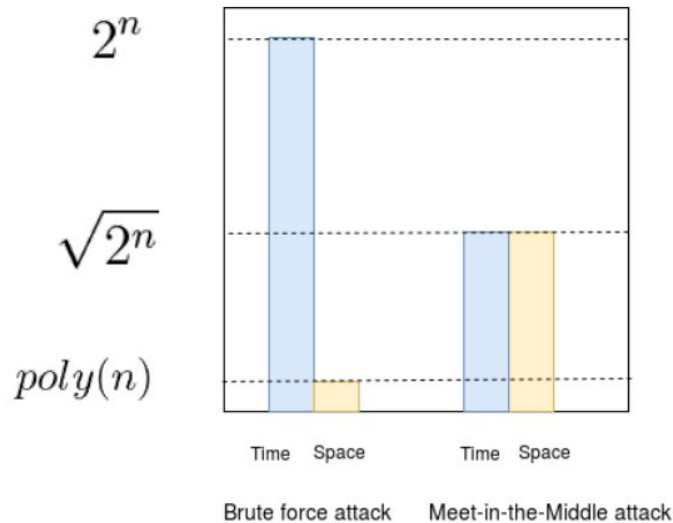
$H = PK = F / G$  (Bit-by-Bit)

Splitting guess for  $G$  into  $G_1, G_2$

$$hG_1 = -hG_2 + f$$

Hash table  $L$

Locality Sensitive hash function



# Physical attacks on how these sparse integers are stored

- A little bit of noise is added to the public key and anything encrypted by AJPS:
  - Public key:  $pk = (R, F \cdot R + G)$
  - Ciphertext:  $(C_1, C_2) = (A \cdot R + B_1, (A \cdot T + B_2) \oplus \mathcal{E}(m))$
- If the attacker can corrupt the storage of this low-Hamming-weight noise, say wipe them out to zero or change them to something the attacker knows, they can recover the secret key  $F$  or the plaintext message  $m$ :
  - $pk = (R, F \cdot R + G) \Rightarrow$  Compute  $R^{-1} \pmod{p}$ , followed by  $R^{-1}(F \cdot R + G) = F$
  - $(C_1, C_2) = (A \cdot R + B_1, (A \cdot T + B_2) \oplus \mathcal{E}(m)) \Rightarrow$  Compute  $R^{-1} \pmod{p}$ . The attacker can then recover  $A = A \cdot R \cdot R^{-1}$ , followed by  $(A \cdot T) \oplus C_2 = (A \cdot T) \oplus (A \cdot T + B_2) \oplus \mathcal{E}(m)$ .
- The job is even easier if we can wipe out  $A$ .
- What if we can make it so that  $B_1 = B_2$ ?

# References

- [1] Divesh Aggarwal, Antoine Joux, Anupam Prakash, and Miklos Santha. A new public-key cryptosystem via Mersenne numbers. Cryptology ePrint Archive, Report 2017/481, version 20171206.004924, 2017.
- [2] Marc Beunardeau, Aisling Connolly, Remi Geraud, and David Naccache. On the hardness of the Mersenne Low Hamming Ratio assumption. Technical report, Cryptology ePrint Archive, 2017/522, 2017.
- [3] Koen de Boer, Leo Ducas, Stacey Jeffery, and Ronald de Wolf. Attacks on the AJPS Mersenne based cryptosystem. Cryptology ePrint Archive, Report 2017/1171, version 20180125.131924, 2018.
- [4] Lov K. Grover. A fast quantum mechanical algorithm for database search. Proceedings of the twenty-eighth annual ACM symposium on Theory of computing, pages 212-219, 1996.